

```
In [1]: #Caden Zadell  
#BSAN 360  
#Lab 4
```

```
In [2]: import pandas as pd  
import numpy as np
```

```
In [3]: #1  
df = pd.read_csv("baseball_hitting.csv")  
df.head()
```

Out[3]:

	Player name	position	Games	At-bat	Runs	Hits	Double (2B)	third baseman	hom runs
0	B Bonds	LF	2986.0	9847.0	2227.0	2935.0	601.0	77.0	762
1	H Aaron	RF	3298.0	12364.0	2174.0	3771.0	624.0	98.0	755
2	B Ruth	RF	2504.0	8399.0	2174.0	2873.0	506.0	136.0	714
3	A Pujols	1B	3080.0	11421.0	1914.0	3384.0	686.0	16.0	703
4	A Rodriguez	SS	2784.0	10566.0	2021.0	3115.0	548.0	31.0	696

```
In [4]: #2  
print("Rows:", df.shape[0])  
print("Columns:", df.shape[1])
```

Rows: 2508  
Columns: 18

```
In [5]: #2  
# This helps confirm that the dataset loaded correctly.  
# The expected dimensions depend on the data source
```

```
In [5]: #3  
print("First five rows:")  
display(df.head())  
  
print("\nLast five rows:")  
display(df.tail())
```

First five rows:

	Player name	position	Games	At-bat	Runs	Hits	Double (2B)	third baseman	home run
0	B Bonds	LF	2986.0	9847.0	2227.0	2935.0	601.0	77.0	762.0
1	H Aaron	RF	3298.0	12364.0	2174.0	3771.0	624.0	98.0	755.0
2	B Ruth	RF	2504.0	8399.0	2174.0	2873.0	506.0	136.0	714.0
3	A Pujols	1B	3080.0	11421.0	1914.0	3384.0	686.0	16.0	703.0
4	A Rodriguez	SS	2784.0	10566.0	2021.0	3115.0	548.0	31.0	696.0

Last five rows:

	Player name	position	Games	At-bat	Runs	Hits	Double (2B)	third baseman	home run
2503	R Derry	LF	187.0	553.0	68.0	124.0	17.0	7.0	
2504	M Kittridge	C	470.0	4027.0	375.0	882.0	108.0	31.0	
2505	G DeMontreville	2B	280.0	3615.0	537.0	1096.0	130.0	35.0	
2506	L Niekro	1B	195.0	499.0	61.0	123.0	26.0	5.0	
2507	B Whitehead	2B	924.0	3316.0	415.0	883.0	100.0	31.0	

```
In [7]: #3
# Confirm that all rows are aligned, headers are correct,
# and that numeric columns are in the right format (no misaligned or m
```

```
In [6]: #4
print(df.columns)

Index(['Player name', 'position', 'Games', 'At-bat', 'Runs', 'Hits',
       'Double (2B)', 'third baseman', 'home run', 'run batted in', 'a
walk',
       'Strikeouts', 'stolen base ', 'Caught stealing', 'AVG',
       'On-base Percentage', 'Slugging Percentage', 'On-base Plus Slugg
ing'],
      dtype='object')
```

```
In [7]: #4
df = df.rename(columns={
    "a walk": "Walked",
    "third baseman": "Triples",
    "home run": "Home Runs",
    "stolen base": "Stolen Bases",
    "run batted in": "RBIs",
```

```
        "Double (2B)": "Doubles",
        "stolen base": "Steals",
        "position": "Position"
    })
```

```
In [8]: print(df.columns)
```

```
Index(['Player name', 'Position', 'Games', 'At-bat', 'Runs', 'Hits', 'Doubles',
       'Triples', 'Home Runs', 'RBIs', 'Walked', 'Strikeouts', 'Steals',
       'Caught stealing', 'AVG', 'On-base Percentage', 'Slugging Percentage',
       'On-base Plus Slugging'],
      dtype='object')
```

```
In [24]: #4
```

```
# Renamed columns for uniformity and to be more clear on what stat is
```

```
In [26]: # 5a. What information is provided in this dataset?
```

```
# Each row represents a player's batting statistics for a given season
# The columns include both basic (Runs, Hits, Home Runs) and advanced
# metrics (On-base Percentage, Slugging Percentage, OPS).
# These stats reflect offensive performance across a variety of metric
```

```
In [9]: print(df.dtypes)
```

Player name	object
Position	object
Games	float64
At-bat	float64
Runs	float64
Hits	float64
Doubles	float64
Triples	float64
Home Runs	float64
RBIs	float64
Walked	float64
Strikeouts	object
Steals	float64
Caught stealing	object
AVG	float64
On-base Percentage	float64
Slugging Percentage	float64
On-base Plus Slugging	float64
dtype: object	

```
In [28]: #5b
```

```
# Most columns are correctly stored as numeric (float64), which is ideal
# 'Strikeouts' and 'Caught stealing' are objects – they should likely
# so these columns may contain commas, text, or missing characters that
# pandas from recognizing them as numbers.
```

```
In [10]: #5c  
df.describe
```

```
Out[10]: <bound method NDFrame.describe of  
mes At-bat Runs Hits Doubles \\  
0 B Bonds LF 2986.0 9847.0 2227.0 2935.0 60  
1.0  
1 H Aaron RF 3298.0 12364.0 2174.0 3771.0 62  
4.0  
2 B Ruth RF 2504.0 8399.0 2174.0 2873.0 50  
6.0  
3 A Pujols 1B 3080.0 11421.0 1914.0 3384.0 68  
6.0  
4 A Rodriguez SS 2784.0 10566.0 2021.0 3115.0 54  
8.0  
... ... ... ... ... ... ... ...  
...  
2503 R Derry LF 187.0 553.0 68.0 124.0 1  
7.0  
2504 M Kittridge C 470.0 4027.0 375.0 882.0 10  
8.0  
2505 G DeMontreville 2B 280.0 3615.0 537.0 1096.0 13  
0.0  
2506 L Niekro 1B 195.0 499.0 61.0 123.0 2  
6.0  
2507 B Whitehead 2B 924.0 3316.0 415.0 883.0 10  
0.0
```

```
Triples Home Runs RBIs Walked Strikeouts Steals Caught st  
ealing \\  
0 77.0 762.0 1996.0 2558.0 1539 514.0  
141  
1 98.0 755.0 2297.0 1402.0 1383 240.0  
73  
2 136.0 714.0 2213.0 2062.0 1330 123.0  
117  
3 16.0 703.0 2218.0 1373.0 1404 117.0  
43  
4 31.0 696.0 2086.0 1338.0 2287 329.0  
76  
... ... ... ... ... ... ...  
...  
2503 7.0 17.0 73.0 78.0 124 2.0  
0  
2504 31.0 17.0 390.0 314.0 166 64.0  
--  
2505 35.0 17.0 497.0 174.0 35 228.0  
--  
2506 5.0 17.0 79.0 29.0 91 0.0  
2  
2507 31.0 17.0 245.0 150.0 138 51.0  
--
```

	Avg	On-base Percentage	Slugging Percentage	On-base Plus Slugging
0	0.298	0.444		0.607
1.051				
1	0.305	0.374		0.555
0.929				
2	0.342	0.474		0.690
1.164				
3	0.296	0.374		0.544
0.918				
4	0.295	0.380		0.550
0.930				
...	...	...		...
...				
2503	0.224	0.322		0.373
0.695				
2504	0.219	0.277		0.274
0.551				
2505	0.303	0.340		0.373
0.713				
2506	0.246	0.288		0.421
0.709				
2507	0.266	0.304		0.331
0.635				

[2508 rows x 18 columns]>

In [11]:

```
#5c
numeric_df = df.select_dtypes(include=["float64", "int64"])
min_values = numeric_df.min()
max_values = numeric_df.max()
mean_values = numeric_df.mean()

print("Minimum values:\n", min_values)
print("\nMaximum values:\n", max_values)
print("\nMean values:\n", mean_values)
```

Minimum values:

Games	2.000
At-bat	262.000
Runs	32.000
Hits	57.000
Doubles	7.000
Triples	0.000
Home Runs	17.000
RBIs	37.000
Walked	19.000
Steals	0.000
AVG	0.123
On-base Percentage	0.157
Slugging Percentage	0.197
On-base Plus Slugging	0.354

dtype: float64

Maximum values:

Games	3562.000
At-bat	14053.000
Runs	2295.000
Hits	4256.000
Doubles	792.000
Triples	309.000
Home Runs	762.000
RBIs	2297.000
Walked	2558.000
Steals	1406.000
AVG	0.367
On-base Percentage	0.482
Slugging Percentage	0.690
On-base Plus Slugging	1.164

dtype: float64

Mean values:

Games	1084.558000
At-bat	3714.962000
Runs	521.644800
Hits	1010.865600
Doubles	181.858000
Triples	32.330800
Home Runs	100.611600
RBIs	494.206400
Walked	373.038000
Steals	76.095200
AVG	0.263320
On-base Percentage	0.331582
Slugging Percentage	0.409925
On-base Plus Slugging	0.741695

dtype: float64

In [35]: #5c

```
# Games: min = 2, max = 3,562
```

```
# Represents very short and very long careers (2 = brief career; 3,562  
#record length, like Pete Rose).  
# Reasonable.  
  
# At-bat: min = 262, max = 14,053  
# Fits realistic career totals (around 10–14k = lifetime leaders such  
#Pete Rose).  
#Reasonable.  
  
# Runs: min = 32, max = 2,295  
# Matches all-time leaders (Rickey Henderson = 2,295).  
# Reasonable.  
  
# Hits: min = 57, max = 4,256  
# 4,256 = Pete Rose's exact hit total, confirming data validity.  
# Reasonable.  
  
# Doubles: min = 7, max = 792  
# 792 ≈ Tris Speaker's career record. Reasonable.  
  
# Triples: min = 0, max = 309  
# 309 ≈ Sam Crawford's record. Reasonable.  
  
# Home Runs: min = 17, max = 762  
# 762 = Barry Bonds' record. Fits perfectly. Reasonable.  
  
# All the data ranges appear reasonable based on the min and max value
```

In [12]:

```
#5d  
print(df.isna().sum())
```

Player name	8
Position	8
Games	8
At-bat	8
Runs	8
Hits	8
Doubles	8
Triples	8
Home Runs	8
RBIs	8
Walked	8
Strikeouts	8
Steals	8
Caught stealing	8
AVG	8
On-base Percentage	8
Slugging Percentage	8
On-base Plus Slugging	20
dtype: int64	

In [44]:

```
#5d  
# Since the value's are missing, we can double check to see if the pla
```

```
# that stat and input a 0.
# If its data entry issue, since the values are low we can delete thos
# players wouldn't be apart of the question we are trying to answer an
```

In [ ]: #6

```
# Research Question 1:
# Who had the best offensive prime (highest AVG, OBP, SLG, OPS)?
# Columns: ['Player name', 'AVG', 'On-base Percentage', 'Slugging Perc

# Research Question 2:
# 2. Which players produced the most runs during their best seasons?
# Columns: ['Player name', 'Home Runs', 'RBIs', 'Runs']
```

In [13]: #7

```
prime_stats = df[['Player name', 'Position', 'AVG', 'On-base Percentage',
                   'Slugging Percentage', 'On-base Plus Slugging',
                   'Home Runs', 'RBIs', 'Hits', 'Steals', 'On-base Plus
prime_stats.head()
```

Out[13]:

	Player name	Position	AVG	On-base Percentage	Slugging Percentage	On-base Plus Slugging	Home Runs	RBIs
0	B Bonds	LF	0.298	0.444	0.607	1.051	762.0	1996.0
1	H Aaron	RF	0.305	0.374	0.555	0.929	755.0	2297.0
2	B Ruth	RF	0.342	0.474	0.690	1.164	714.0	2213.0
3	A Pujols	1B	0.296	0.374	0.544	0.918	703.0	2218.0
4	A Rodriguez	SS	0.295	0.380	0.550	0.930	696.0	2086.0

In [ ]: #7

```
# This DataFrame isolates the core stats that define a player's offense
# It will be used to evaluate players using summary statistics, scoring
```

In [ ]: #8

```
# Who had the best offensive prime (highest AVG, OBP, SLG, OPS)?
# Method: Ranking analysis
# Compute top 10 players by On-base Plus Slugging (OPS)
# Calculate z-scores or percentiles to normalize across eras.

# Research Question 2:
# Which players generated the most runs during their peak?
# Method: Correlation and ranking
# Use correlation matrix to test relationship between HRs, RBIs, Runs,
# Rank players based on combined weighted score.
```

```
In [ ]: #Caden Zadell  
#Project 2 Assignment Work Starts Here
```

```
In [7]: #Data Cleaning and Prep
```

```
In [ ]:
```

```
In [14]: print(df.head())
```

	Player name	Position	Games	At-bat	Runs	Hits	Doubles	Trip
0	B Bonds	LF	2986.0	9847.0	2227.0	2935.0	601.0	7
1	H Aaron	RF	3298.0	12364.0	2174.0	3771.0	624.0	9
2	B Ruth	RF	2504.0	8399.0	2174.0	2873.0	506.0	13
3	A Pujols	1B	3080.0	11421.0	1914.0	3384.0	686.0	1
4	A Rodriguez	SS	2784.0	10566.0	2021.0	3115.0	548.0	3

	Home Runs	RBI's	Walked	Strikeouts	Steals	Caught stealing	AVG
0	762.0	1996.0	2558.0	1539	514.0	141	0.298
1	755.0	2297.0	1402.0	1383	240.0	73	0.305
2	714.0	2213.0	2062.0	1330	123.0	117	0.342
3	703.0	2218.0	1373.0	1404	117.0	43	0.296
4	696.0	2086.0	1338.0	2287	329.0	76	0.295

	On-base Percentage	Slugging Percentage	On-base Plus Slugging
0	0.444	0.607	1.051
1	0.374	0.555	0.929
2	0.474	0.690	1.164
3	0.374	0.544	0.918
4	0.380	0.550	0.930

```
In [15]: df_clean = df[['Player name', 'AVG', 'On-base Percentage', 'Slugging Percentage',  
'On-base Plus Slugging', 'Home Runs', 'RBI's', 'Runs', '']
```

```
In [16]: print(df_clean.isnull().sum())
```

```
Player name      8
AVG             8
On-base Percentage  8
Slugging Percentage  8
On-base Plus Slugging 20
Home Runs        8
RBIs            8
Runs             8
Hits             8
Steals           8
dtype: int64
```

```
In [17]: print(df_clean.dtypes)
```

```
Player name      object
AVG             float64
On-base Percentage  float64
Slugging Percentage  float64
On-base Plus Slugging  float64
Home Runs        float64
RBIs            float64
Runs             float64
Hits             float64
Steals           float64
dtype: object
```

```
In [38]: #Missing Data
```

```
In [18]: df_clean = df_clean.dropna(subset=['AVG', 'On-base Percentage', 'Slugg
df_clean['Steals'] = df_clean['Steals'].fillna(0)

print("Missing values:")
print(df_clean.isnull().sum())

print("\nRows after cleaning:", len(df_clean))
```

Missing values:

```
Player name      0
AVG             0
On-base Percentage  0
Slugging Percentage  0
On-base Plus Slugging  0
Home Runs        0
RBIs            0
Runs             0
Hits             0
Steals           0
dtype: int64
```

Rows after cleaning: 2488

```
In [19]: print(df_clean.shape)
```

```
(2488, 10)
```

```
In [42]: #Data Transformation
```

```
In [20]: from sklearn.preprocessing import StandardScaler
```

```
In [21]: numeric_cols = ['AVG', 'On-base Percentage', 'Slugging Percentage',  
                     'On-base Plus Slugging', 'Home Runs', 'RBIs', 'Runs',
```

```
In [22]: scaler = StandardScaler()
```

```
In [23]: scaled_data = scaler.fit_transform(df_clean[numeric_cols])  
scaled_df = pd.DataFrame(scaled_data, columns=numeric_cols)
```

```
In [24]: scaled_df['Player name'] = df_clean['Player name'].values
```

```
In [25]: print(scaled_df.head())
```

	AVG	On-base Percentage	Slugging Percentage	On-base Plus	Slug
0	1.402200 9191	3.660724	3.933776		4.29
1	1.684847 3450	1.381075	2.894876		2.60
2	3.178838 9837	4.637717	5.592020		5.86
3	1.321443 0555	1.381075	2.675109		2.45
4	1.281065 7350	1.576474	2.794982		2.61

	Home Runs	RBIs	Runs	Hits	Steals	Player name
0	6.604641	4.133466	4.481532	2.819983	3.852164	B Bonds
1	6.534704	4.961928	4.342289	4.044516	1.444671	H Aaron
2	6.125074	4.730729	4.342289	2.729168	0.416654	B Ruth
3	6.015173	4.744491	3.659207	3.477657	0.363936	A Pujols
4	5.945236	4.381179	3.940322	3.083638	2.226667	A Rodriguez

```
In [26]: print(scaled_df[numeric_cols].mean())  
print(scaled_df[numeric_cols].std())
```

```
AVG           -3.655525e-16
On-base Percentage   1.005269e-15
Slugging Percentage  3.655525e-16
On-base Plus Slugging 6.854110e-16
Home Runs          -9.138813e-17
RBIs              0.000000e+00
Runs               -1.827763e-16
Hits               -4.569407e-17
Steals             -1.142352e-17
dtype: float64
AVG           1.000201
On-base Percentage   1.000201
Slugging Percentage  1.000201
On-base Plus Slugging 1.000201
Home Runs          1.000201
RBIs              1.000201
Runs               1.000201
Hits               1.000201
Steals             1.000201
dtype: float64
```

```
In [27]: #Caden Zadell
#Project 3 Assignment Starts Here
import pandas as pd

df = pd.read_csv("baseball_hitting.csv")

df.columns = df.columns.str.strip()

df = df.rename(columns={
    "position": "Position",
    "Double (2B)": "Doubles",
    "third baseman": "Triples",
    "home run": "Home Runs",
    "run batted in": "RBIs",
    "a walk": "Walked",
    "stolen base ": "Steals",    #
    "stolen base": "Steals"
})

df = df.dropna(subset=["Player name", "AVG", "On-base Percentage", "Slugging"])

df["On-base Plus Slugging"] = df["On-base Percentage"] + df["Slugging"]

df = df[df["Games"] >= 30]

df.to_csv("cleaned_baseball_data.csv", index=False)

print("Cleaned data shape:", df.shape)
print(df.columns)
```

```
Cleaned data shape: (2496, 18)
Index(['Player name', 'Position', 'Games', 'At-bat', 'Runs', 'Hits', 'Doubles',
       'Triples', 'Home Runs', 'RBIs', 'Walked', 'Strikeouts', 'Steals',
       'Caught stealing', 'AVG', 'On-base Percentage', 'Slugging Percentage',
       'On-base Plus Slugging'],
      dtype='object')
```

In [28]: #Professor Feedback

```
df = df.dropna(subset=["Player name"])

df["On-base Plus Slugging"] = df["On-base Percentage"] + df["Slugging"]

print(df["On-base Plus Slugging"].isna().sum())

df = df[df["Games"] >= 30]

print(len(df))
```

```
0
2496
```

In [29]: #Processing Strings

```
df.columns = df.columns.str.strip()

df["Player name"] = df["Player name"].str.strip()

print(df["Player name"].head())
```

```
0      B Bonds
1      H Aaron
2      B Ruth
3      A Pujols
4      A Rodriguez
Name: Player name, dtype: object
```

In [ ]: #Combining and Merging Datasets

```
#Only have one dataset for this project.
#No merging or combining is needed
```

In [31]: #Reshaping and Pivoting

```
pivot_position = df.pivot_table(
    values=["AVG", "On-base Percentage", "Slugging Percentage", "On-base Plus Slugging"],
    index="Position",
    aggfunc="mean"
).round(3)

print("Average Offensive Performance by Position:")
print(pivot_position)
```

Average Offensive Performance by Position:			
Position	AVG	On-base Percentage	On-base Plus Slugging \ Slugging Percentage
1B	0.268	0.342	0.777
2B	0.263	0.327	0.711
3B	0.262	0.328	0.736
C	0.249	0.319	0.706
CF	0.267	0.335	0.744
DH	0.270	0.350	0.810
LF	0.269	0.338	0.767
OF	0.268	0.336	0.751
P	0.221	0.279	0.619
RF	0.269	0.339	0.771
SS	0.260	0.319	0.697

Position	Slugging Percentage
1B	0.435
2B	0.383
3B	0.408
C	0.387
CF	0.409
DH	0.460
LF	0.428
OF	0.415
P	0.340
RF	0.432
SS	0.378

```
In [32]: df_clean = df.copy()
df_clean.to_csv("cleaned_baseball_data.csv", index=False)
```

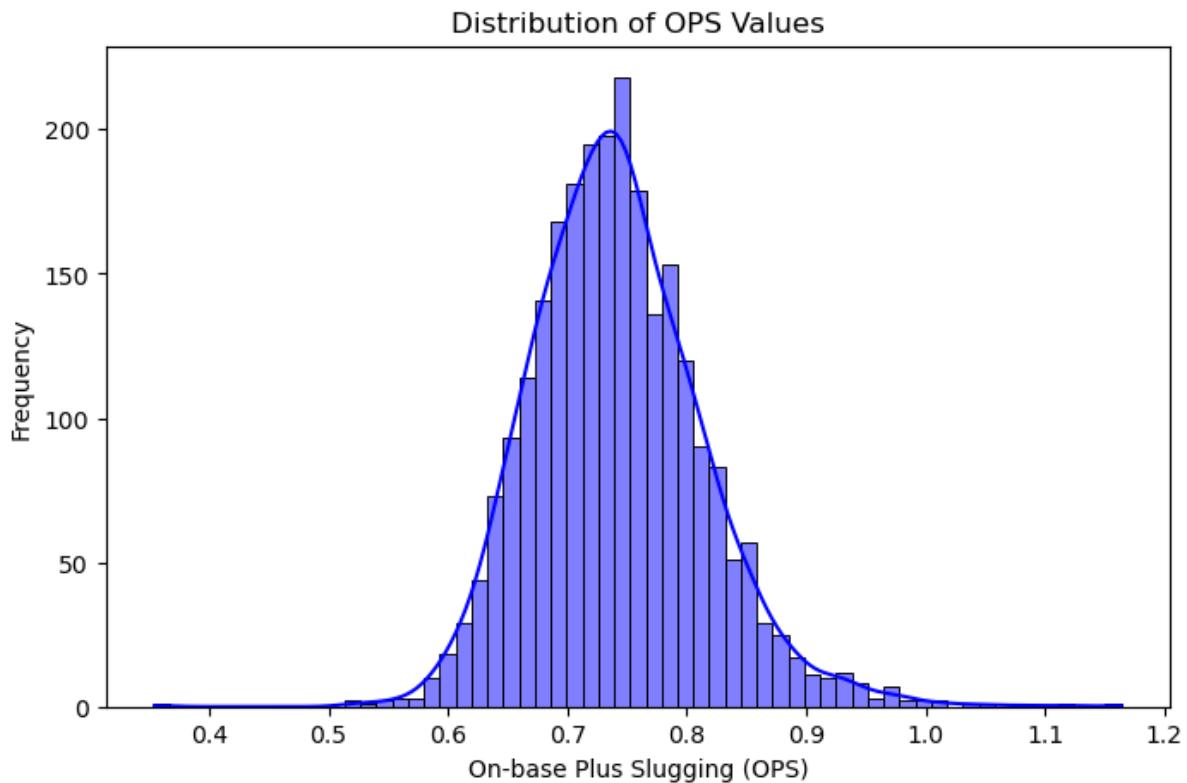
```
In [1]: #Caden Zadell
#BSAN
#Project 4 Assignment Starts Here
```

```
In [33]: import matplotlib.pyplot as plt
import seaborn as sns

df = pd.read_csv("cleaned_baseball_data.csv")
```

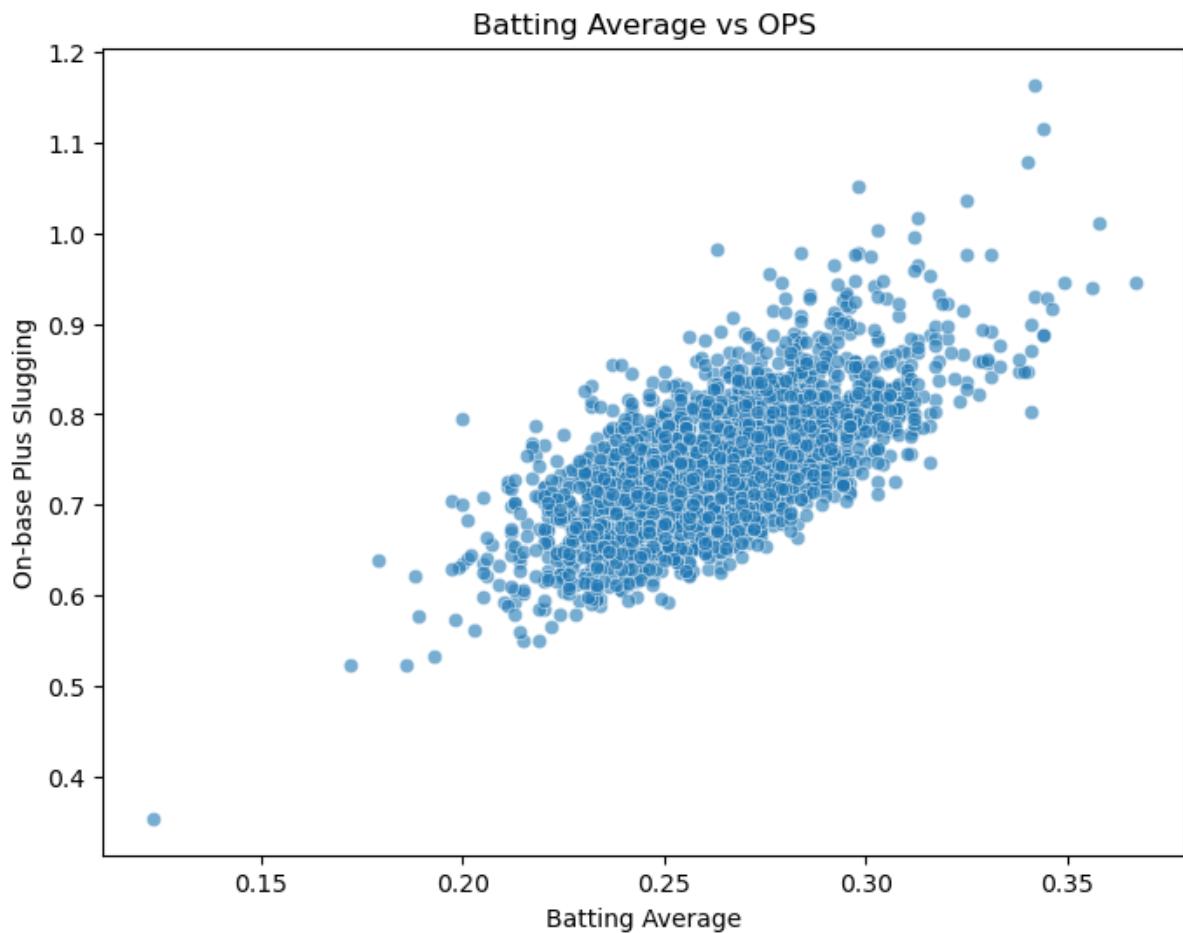
```
In [34]: #1 Distribution of OPS

plt.figure(figsize=(8, 5))
sns.histplot(df["On-base Plus Slugging"], kde=True, color="blue")
plt.title("Distribution of OPS Values")
plt.xlabel("On-base Plus Slugging (OPS)")
plt.ylabel("Frequency")
plt.show()
```



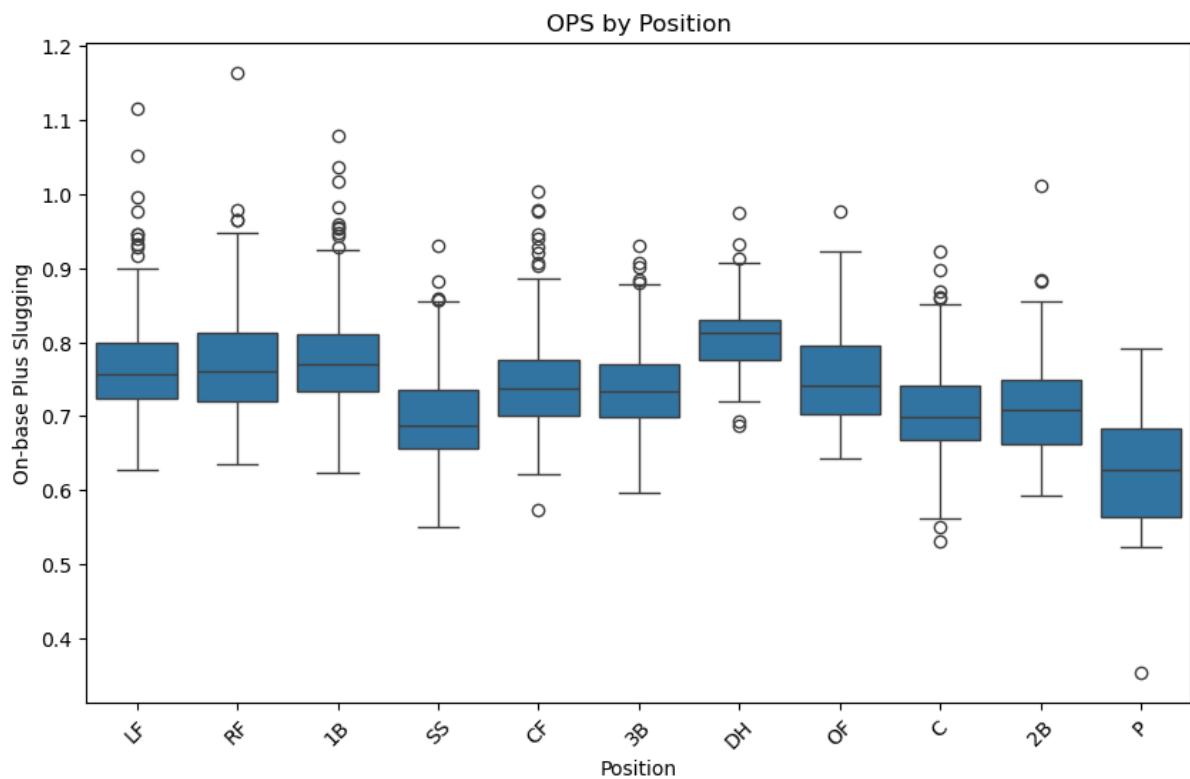
In [35]: #2 AVG vs OPS Scatter Plot

```
plt.figure(figsize=(8, 6))
sns.scatterplot(data=df, x="AVG", y="On-base Plus Slugging", alpha=0.6
plt.title("Batting Average vs OPS")
plt.xlabel("Batting Average")
plt.ylabel("On-base Plus Slugging")
plt.show()
```



In [36]: # 3. Boxplot of OPS by Position

```
plt.figure(figsize=(10, 6))
sns.boxplot(data=df, x="Position", y="On-base Plus Slugging")
plt.xticks(rotation=45)
plt.title("OPS by Position")
plt.xlabel("Position")
plt.ylabel("On-base Plus Slugging")
plt.show()
```



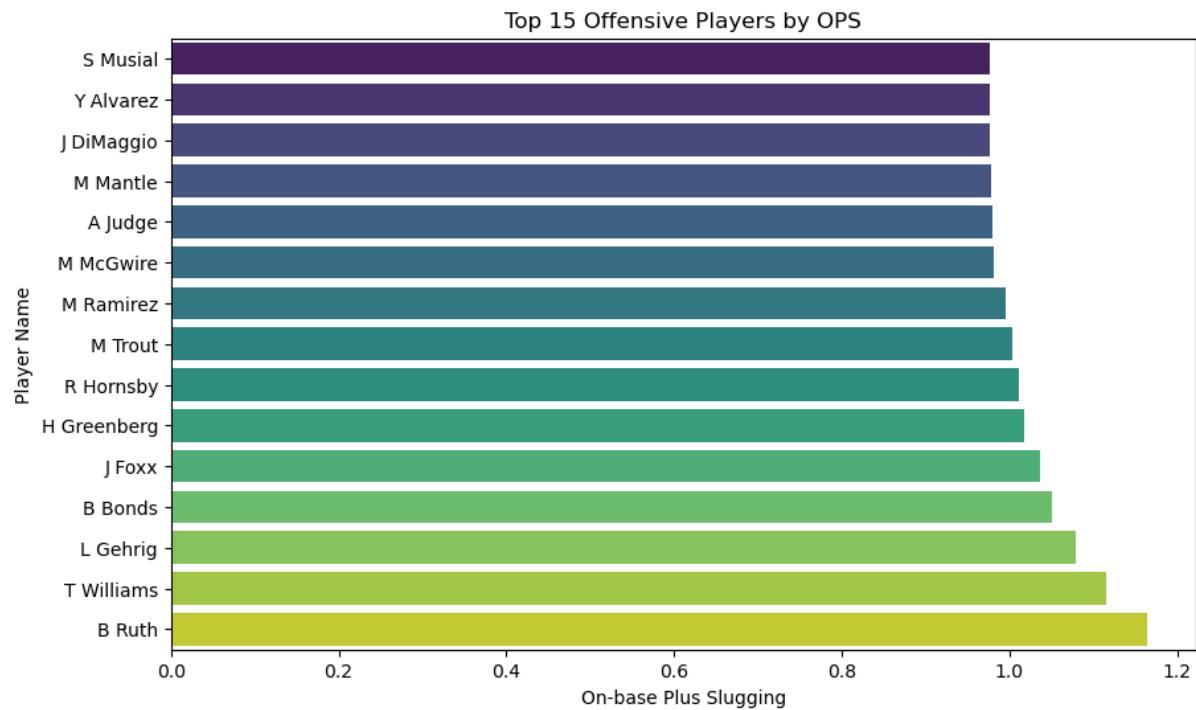
In [37]: #4 Top 15 Players by OPS

```
top_ops = df.nlargest(15, "On-base Plus Slugging")[['Player name', "On  
plt.figure(figsize=(10, 6))  
sns.barplot(data=top_ops, x="On-base Plus Slugging", y="Player name",  
plt.title("Top 15 Offensive Players by OPS")  
plt.xlabel("On-base Plus Slugging")  
plt.ylabel("Player Name")  
plt.show()
```

/var/folders/x9/5gqcqk055lb1bfdx10rnvrcw0000gn/T/ipykernel\_3177/2071071  
647.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=top_ops, x="On-base Plus Slugging", y="Player name",  
palette="viridis")
```



```
In [38]: #Caden Zadell
#BSAN 360
#Project 5 Assignment
import pandas as pd

df = pd.read_csv("cleaned_baseball_data.csv")
print(df.head())
```

	Player name	Position	Games	At-bat	Runs	Hits	Doubles	Trip
0	B Bonds	LF	2986.0	9847.0	2227.0	2935.0	601.0	7
1	H Aaron	RF	3298.0	12364.0	2174.0	3771.0	624.0	9
2	B Ruth	RF	2504.0	8399.0	2174.0	2873.0	506.0	13
3	A Pujols	1B	3080.0	11421.0	1914.0	3384.0	686.0	1
4	A Rodriguez	SS	2784.0	10566.0	2021.0	3115.0	548.0	3

	Home Runs	RBI's	Walked	Strikeouts	Steals	Caught stealing	AVG
0	762.0	1996.0	2558.0	1539	514.0	141	0.298
1	755.0	2297.0	1402.0	1383	240.0	73	0.305
2	714.0	2213.0	2062.0	1330	123.0	117	0.342
3	703.0	2218.0	1373.0	1404	117.0	43	0.296
4	696.0	2086.0	1338.0	2287	329.0	76	0.295

	On-base Percentage	Slugging Percentage	On-base Plus Slugging
0	0.444	0.607	1.051
1	0.374	0.555	0.929
2	0.474	0.690	1.164
3	0.374	0.544	0.918
4	0.380	0.550	0.930

In [2]: #New Research Question  
*#Do naturally occurring clusters of MLB players reveal different offensive styles (power hitters, contact hitters, balanced hitters), and which statistics most define each group?*

In [3]: #Data Aggregation

```
In [39]: numeric_cols = [
    "AVG", "On-base Percentage", "Slugging Percentage",
    "On-base Plus Slugging", "Home Runs", "RBI's",
    "Runs", "Hits", "Steals"
]

overall_summary = df[numeric_cols].agg(["mean", "median", "min", "max"])

print("Overall summary of key offensive stats:")
print(overall_summary)
```

Overall summary of key offensive stats:

	AVG	On-base Percentage	Slugging Percentage	On-base Plus Sl	
ugging \					
mean	0.263	0.332		0.410	
	0.741				
median	0.262	0.330		0.407	
	0.737				
min	0.123	0.157		0.197	
	0.354				
max	0.367	0.482		0.690	
	1.164				
std	0.025	0.031		0.050	
	0.072				
	Home Runs	RBIs	Runs	Hits	Steals
mean	100.698	493.929	520.849	1010.128	75.701
median	69.000	404.000	423.000	853.000	32.000
min	17.000	37.000	32.000	57.000	0.000
max	762.000	2297.000	2295.000	4256.000	1406.000
std	100.032	362.475	379.204	681.278	113.160

In [40]: #Data grouping by position

```
position_summary = (
    df.groupby("Position")
    .agg({
        "On-base Plus Slugging": ["mean", "median", "max"],
        "Home Runs": ["mean", "max"],
        "RBIs": "mean",
        "Player name": "count"
    })
    .round(3)
)
```

In [41]: position\_summary = position\_summary.rename(columns={"Player name": "Pl

```
print("Offensive summary by position:")
print(position_summary)
```

Position	On-base Plus Slugging			Home Runs			RBIs
		mean	median	max	mean	max	mean
	1B	0.777	0.771	1.079	131.099	703.0	589.958
2B		0.711	0.709	1.011	70.487	377.0	448.857
3B		0.736	0.734	0.930	102.859	548.0	501.757
C		0.706	0.698	0.922	76.521	427.0	366.743
CF		0.744	0.738	1.003	97.263	660.0	488.737
DH		0.810	0.813	0.974	193.486	541.0	710.286
LF		0.767	0.756	1.116	109.405	762.0	497.060
OF		0.751	0.741	0.976	90.947	475.0	488.120
P		0.619	0.627	0.791	26.786	61.0	169.143
RF		0.771	0.761	1.164	122.700	755.0	538.672
SS		0.697	0.688	0.930	78.978	696.0	510.780

Position	Player count	
	count	count
1B	354	
2B	273	
3B	284	
C	338	
CF	274	
DH	35	
LF	333	
OF	75	
P	14	
RF	293	
SS	223	

In [9]: #Grouping by OPS

```
In [42]: ops_bins = [0.0, 0.700, 0.800, 0.900, 1.200]
ops_labels = [
    "Below Avg (<.700)",
    "Above Avg (.700-.800)",
    "Great (.800-.900)",
    "Elite (>.900)"
]

df["OPS_tier"] = pd.cut(
    df["On-base Plus Slugging"],
    bins=ops_bins,
    labels=ops_labels,
    include_lowest=True
)

tier_counts = df.groupby("OPS_tier")["Player name"].count()
```

```
/var/folders/x9/5gqcqk055lb1bfdx10rnvrcw0000gn/T/ipykernel_3177/2228173  
739.py:16: FutureWarning: The default of observed=False is deprecated a  
nd will be changed to True in a future version of pandas. Pass observed  
=False to retain current behavior or observed=True to adopt the future  
default and silence this warning.
```

```
tier_counts = df.groupby("OPS_tier")["Player name"].count()
```

```
In [43]: print("\nNumber of players in each OPS tier:")  
print(tier_counts)
```

```
Number of players in each OPS tier:  
OPS_tier  
Below Avg (<.700)      714  
Above Avg (.700-.800)   1320  
Great (.800-.900)     402  
Elite (>.900)         60  
Name: Player name, dtype: int64
```

```
In [13]: #Pivot Table
```

```
In [44]: pivot_ops_mean = pd.pivot_table(  
        df,  
        values="On-base Plus Slugging",  
        index="Position",  
        columns="OPS_tier",  
        aggfunc="mean"  
    ).round(3)
```

```
/var/folders/x9/5gqcqk055lb1bfdx10rnvrcw0000gn/T/ipykernel_3177/1064842  
689.py:1: FutureWarning: The default value of observed=False is depreca  
ted and will change to observed=True in a future version of pandas. Spe  
cify observed=False to silence this warning and retain the current beha  
vior
```

```
pivot_ops_mean = pd.pivot_table(
```

```
In [45]: print("Average OPS by position and OPS tier:")  
print(pivot_ops_mean)
```

Average OPS by position and OPS tier:

OPS\_tier Below Avg (<.700) Above Avg (.700-.800) Great (.800-.900)

\

Position

1B	0.674	0.754	0.840
2B	0.657	0.744	0.837
3B	0.667	0.744	0.831
C	0.660	0.741	0.830
CF	0.668	0.746	0.835
DH	0.690	0.768	0.828
LF	0.678	0.749	0.835
OF	0.676	0.745	0.825
P	0.592	0.781	NaN
RF	0.672	0.746	0.832
SS	0.654	0.740	0.841

OPS\_tier Elite (>.900)

Position

1B	0.959
2B	1.011
3B	0.913
C	0.922
CF	0.945
DH	0.932
LF	0.975
OF	0.950
P	NaN
RF	0.949
SS	0.930

```
In [46]: ops_position_crosstab = pd.crosstab(df["Position"], df["OPS_tier"])
```

```
print("\nCross-tab of number of players by position and OPS tier:")
print(ops_position_crosstab)
```

Cross-tab of number of players by position and OPS tier:

OPS_tier	Below Avg (<.700)	Above Avg (.700-.800)	Great (.800-.900)
\			
Position			
1B	38	208	93
2B	127	127	18
3B	73	176	32
C	177	134	26
CF	69	162	34
DH	2	13	16
LF	37	213	73
OF	18	39	16
P	12	2	0
RF	35	162	82
SS	126	84	12

OPS\_tier Elite (>.900)

Position	
1B	15
2B	1
3B	3
C	1
CF	9
DH	4
LF	10
OF	2
P	0
RF	14
SS	1

```
In [47]: #Caden Zadell
#Project Assignment 6
import pandas as pd

df = pd.read_csv("cleaned_baseball_data.csv")
print(df.head())
print(df.columns)
```

```

      Player name Position  Games  At-bat   Runs   Hits  Doubles  Trip
les \
0     B Bonds        LF  2986.0  9847.0  2227.0  2935.0   601.0    7
7.0
1     H Aaron        RF  3298.0 12364.0  2174.0  3771.0   624.0    9
8.0
2     B Ruth         RF  2504.0  8399.0  2174.0  2873.0   506.0   13
6.0
3     A Pujols       1B  3080.0 11421.0  1914.0  3384.0   686.0    1
6.0
4     A Rodriguez    SS  2784.0 10566.0  2021.0  3115.0   548.0    3
1.0

      Home Runs    RBIs  Walked Strikeouts  Steals Caught stealing    AVG
\
0     762.0  1996.0  2558.0       1539  514.0           141  0.298
1     755.0  2297.0  1402.0       1383  240.0            73  0.305
2     714.0  2213.0  2062.0       1330  123.0           117  0.342
3     703.0  2218.0  1373.0       1404  117.0            43  0.296
4     696.0  2086.0  1338.0       2287  329.0            76  0.295

      On-base Percentage  Slugging Percentage  On-base Plus Slugging
0             0.444                  0.607          1.051
1             0.374                  0.555          0.929
2             0.474                  0.690          1.164
3             0.374                  0.544          0.918
4             0.380                  0.550          0.930
Index(['Player name', 'Position', 'Games', 'At-bat', 'Runs', 'Hits', 'Doubles',
       'Triples', 'Home Runs', 'RBIs', 'Walked', 'Strikeouts', 'Steals',
       'Caught stealing', 'AVG', 'On-base Percentage', 'Slugging Percentage',
       'On-base Plus Slugging'],
      dtype='object')

```

In [2]: #K Means Clustering of Player Archetypes

```

In [48]: from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

cluster_features = [
    "AVG", "On-base Percentage", "Slugging Percentage",
    "Home Runs", "RBIs", "Runs", "Hits", "Steals"
]

X = df[cluster_features].dropna()

scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

kmeans = KMeans(n_clusters=3, random_state=42)

```

```
clusters = kmeans.fit_predict(X_scaled)

X["Cluster"] = clusters
df["Cluster"] = clusters
```

```
In [50]: cluster_profiles = df.groupby("Cluster")[cluster_features].mean().round(2)
cluster_profiles
```

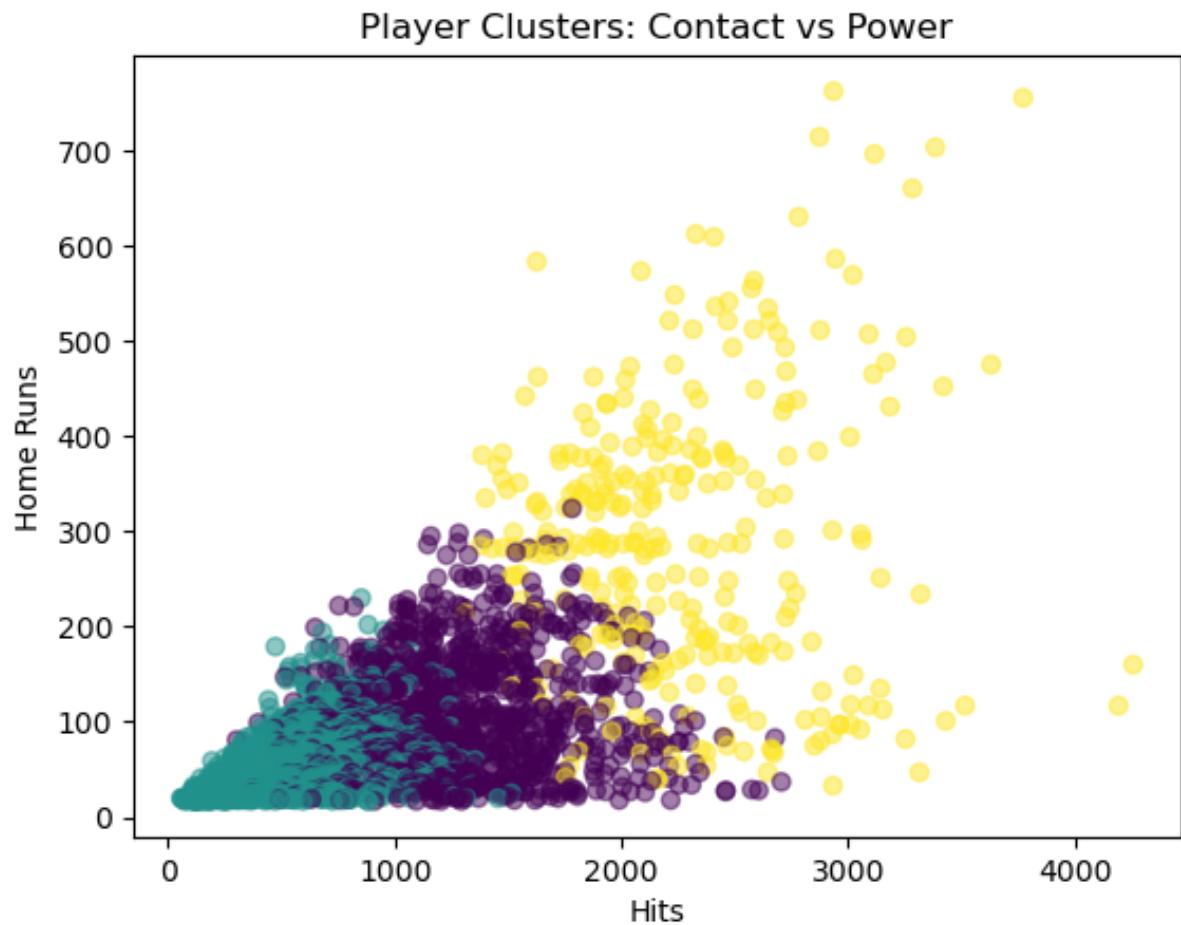
```
Out[50]:
```

	AVG	On-base Percentage	Slugging Percentage	Home Runs	RBIs	Runs	Hits	Steal
Cluster								
0	0.28	0.34	0.42	106.75	592.80	655.40	1276.82	105.9
1	0.25	0.31	0.39	52.21	244.40	247.39	513.53	26.4
2	0.29	0.37	0.48	289.02	1242.66	1256.08	2272.49	188.1

```
In [ ]:
```

```
In [51]: import matplotlib.pyplot as plt

plt.scatter(df["Hits"], df["Home Runs"], c=df["Cluster"], alpha=0.5)
plt.xlabel("Hits")
plt.ylabel("Home Runs")
plt.title("Player Clusters: Contact vs Power")
plt.show()
```



```
In [ ]:
```

```
In [52]: cluster_profiles.std(axis=0).sort_values(ascending=False)
```

```
Out[52]: Hits          882.034642
          Runs         507.402559
          RBIs         506.659584
          Home Runs    124.013400
          Steals        81.195589
          Slugging Percentage  0.045826
          On-base Percentage  0.030000
          AVG           0.020817
          dtype: float64
```

```
In [ ]:
```